

HPG

FACE RECOGNITION

Using PCA, LDA, ICA, Fisherfaces



Rohan Vinkare

S20220010244

INTRODUCTION

- Face recognition is a widely adopted biometric technique used in surveillance, security, and identity verification systems.
- This project evaluates and compares the performance of four classical facial recognition algorithms: Principal Component Analysis (PCA), Linear Discriminant Analysis (LDA), Independent Component Analysis (ICA), and Fisherfaces.
- The focus is on preparing these algorithms for High Performance Computing (HPC) optimization using parallel computing platforms such as CUDA and OpenMP.
- As a baseline, sequential (serial) implementations of all four algorithms have been successfully completed and evaluated.

OBJECTIVE

- Implement and analyze four classical face recognition algorithms: PCA, LDA, ICA, and Fisherfaces, using a common dataset for consistent evaluation.
- Evaluate their performance in sequential form to establish a baseline for comparison with parallel implementations.
- The project aims to optimize computational performance using parallel processing techniques tailored to each algorithm.
- Key performance metrics such as execution time, speedup, and scalability will be measured to assess the effectiveness of each parallel approach.

DATASET AND TOOLS

- Dataset: AT&T Face Dataset on Kaggle
- 40 subjects, 10 grayscale images each (size: 112x92).

Environment :

- Multi-core CPU (for OpenMP)
- CUDA-compatible GPU (for parallel acceleration)

Tools :

- C++ for algorithm logic
- openCV
- CUDA/OpenMP (planned)

- PCA (Principal Component Analysis):

Successfully implemented for face recognition.

Used to reduce dimensionality and extract key features (eigenfaces).

- LDA (Linear Discriminant Analysis):

Implemented in sequential form to enhance class separability.

- ICA (Independent Component Analysis):

For feature extraction based on statistical independence.

- Fisherfaces (PCA + LDA):

Combined approach to improve recognition accuracy over PCA alone.

**WORK
DONE SO
FAR**

METHODOLOGY OVERVIEW

Algorithm	Parallel Platform (Planned)	Reason
PCA	CUDA	Matrix operations & eigen decomposition benefit from GPU
Fisherfaces	CUDA	Combines PCA & LDA → GPU ideal for high-dimensional ops
LDA	OpenMP	Loop-level parallelism in scatter matrix computation
ICA	OpenMP	Independent component stages map well to shared memory threads

WHY THESE PLATFORMS?

- **CUDA (for PCA & Fisherfaces):**
 - High throughput for matrix ops
 - Significant speedup on GPUs
- **OpenMP (for LDA & ICA):**
 - Easy integration with C++
 - Best for loop-heavy or memory-shared tasks

WHY NOT OTHERS?

- **MPI:** More suitable for distributed clusters
- **Pthreads:** Lower-level, harder to manage than OpenMP

PCA – SUMMARY

- Purpose: Reduce dimensionality via eigenfaces.
- **Steps:**
 - Normalize faces
 - Compute covariance matrix
 - Perform eigen decomposition
- Current: Implemented in C++ (sequential)
- Planned: CUDA for speedup via parallel projection & eigenspace computation

LDA – SUMMARY

- Purpose: Maximize class separability.
- Requires: Supervised learning (label information).
- **Steps:**
 - Compute between-class & within-class scatter matrices
 - Solve generalized eigenproblem
- Planned: OpenMP for parallel scatter matrix computation

ICA – SUMMARY

- Purpose: Extract independent statistical features.
- **Steps:**
 - Whitening
 - Iterative estimation of independent components
- **Planned: OpenMP**
 - Each component estimation as parallel loop
 - Shared memory efficiency

FISHERFACES – SUMMARY

- Purpose: Combine PCA + LDA for optimal class separation.
- **Steps:**
 - PCA for dimensionality reduction
 - LDA for maximizing discriminability
- **Planned:** CUDA (both PCA and LDA stages are compute-heavy)

SAMPLE RESULTS :- OPENMP

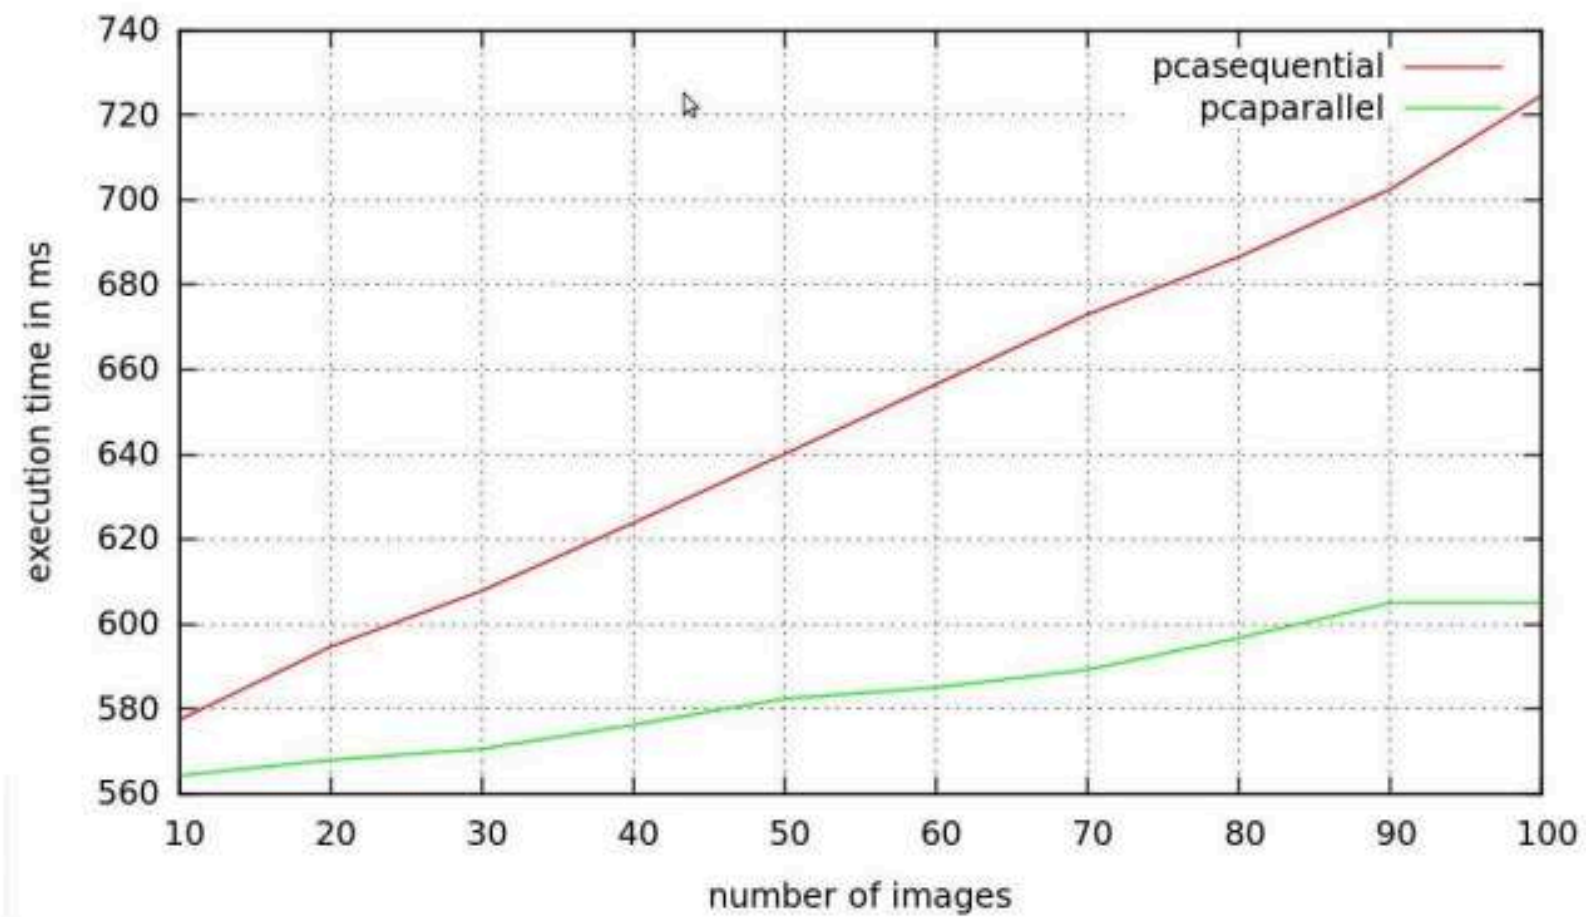


Fig. 2. Performance comparison of sequential and parallel PCA Algorithm.




Fig. 3. Performance comparison of sequential and parallel ICA.

SAMPLE RESULTS :- OPENMP




Fig.4. Performance Comparison of LDA sequential and parallel algorithms.


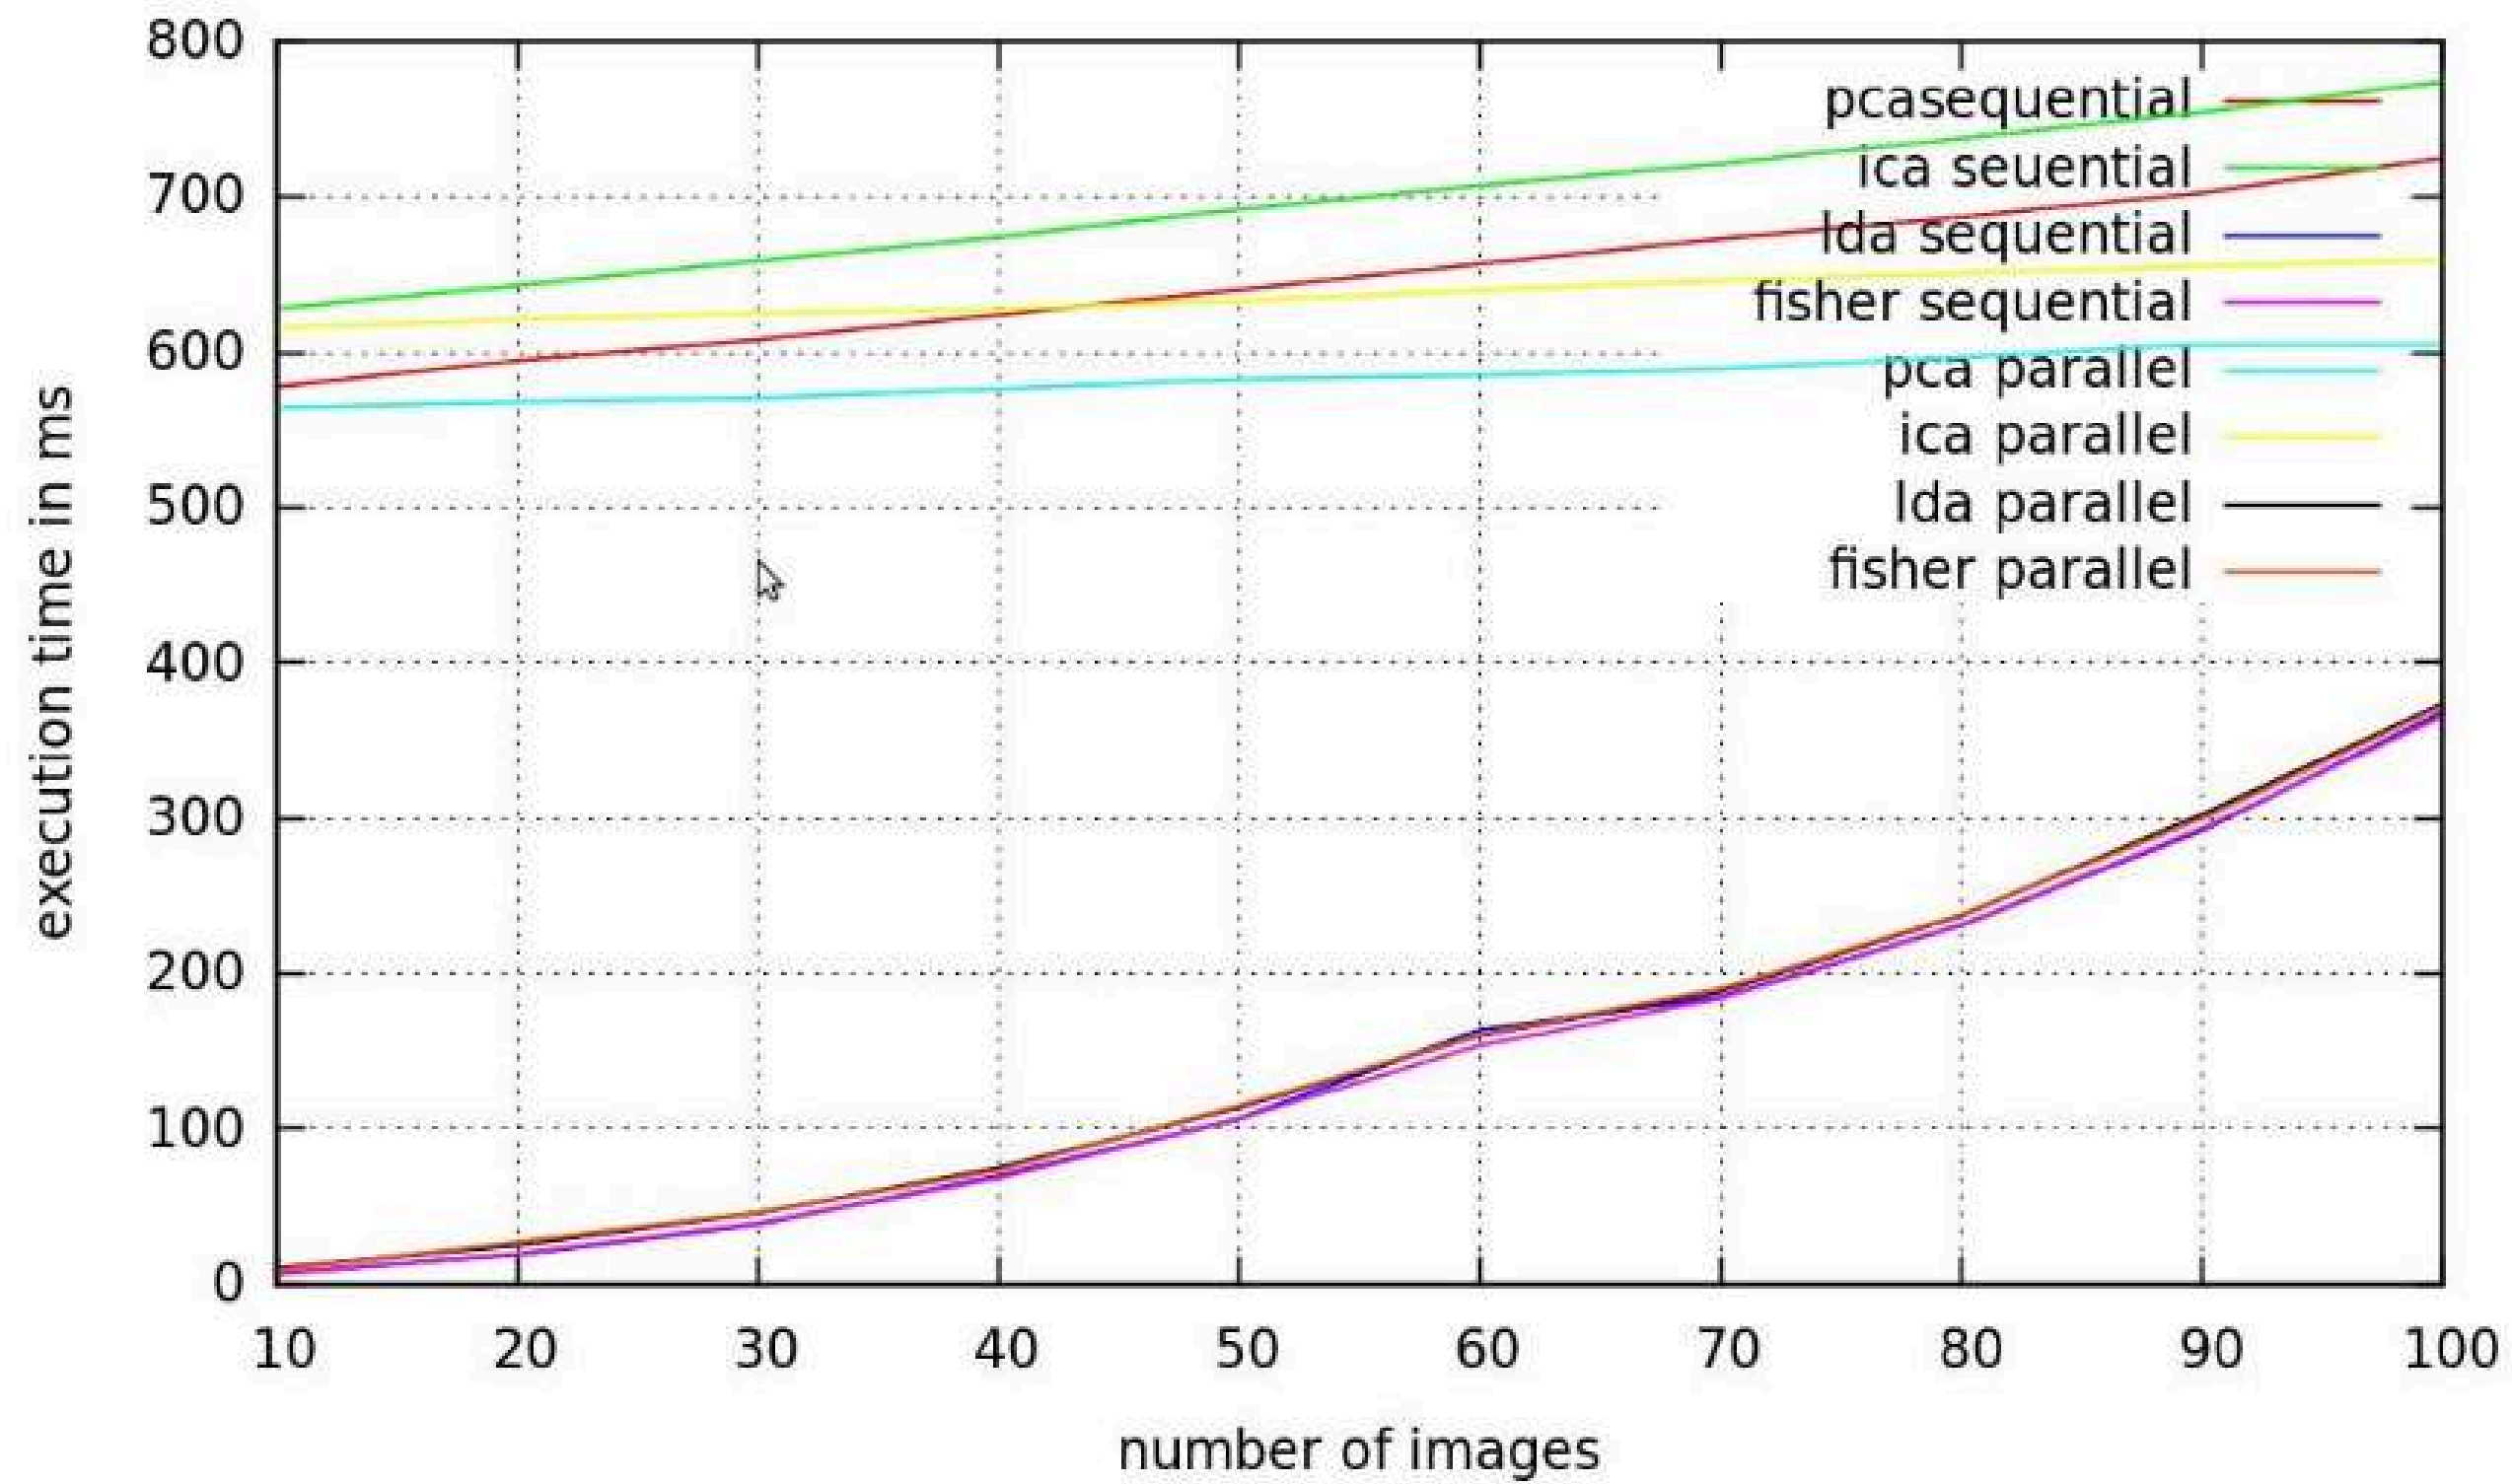


Fig. 5. Performance Comparison of sequential and parallel Fisher Face.



**SAMPLE
RESULTS
OPENMP**

FUTURE WORK

- **PCA** will be parallelized using CUDA to leverage GPU acceleration and improve performance for large-scale datasets.
- **LDA** will be parallelized using OpenMP, enabling multi-threaded execution on multi-core CPUs for improved efficiency.
- **Benchmark Performance**
 - Measure execution time, speedup, and scalability for each parallel implementation.
 - Compare against sequential baselines to evaluate improvement.

Evaluation - 02

HPG

FACE RECOGNITION

Using PCA, LDA, Fisherfaces



Guide :- Dr. Bheemappa Halavar

Sujal Awargand
S20220010213

Rohan Vinkare
S20220010244

Jathin Reddy
S20220010241

PROJECT OVERVIEW

PROBLEM STATEMENT

- Facial recognition systems require significant computational resources due to high-dimensional data processing
- Traditional sequential implementations of PCA and LDA algorithms are computationally intensive
- Need for faster processing for real-time facial recognition applications
- Challenge of optimizing performance while maintaining accuracy

OBJECTIVES

1. Implement and optimize PCA and LDA algorithms for facial recognition
2. Develop parallel implementations using:
 - OpenMP (CPU parallel processing)
 - CUDA (GPU acceleration)
3. Compare performance metrics between:
 - Sequential implementation
 - OpenMP parallel implementation
 - CUDA parallel implementation
4. Achieve significant speedup while maintaining recognition accuracy

DATASET AND TOOLS

- Dataset: AT&T Face Dataset on Kaggle
- 40 subjects, 10 grayscale images each (size: 112x92).

Environment :

- Multi-core CPU (for OpenMP)
- CUDA-compatible GPU (for parallel acceleration)

Tools :

- C++ for algorithm logic
- openCV
- CUDA/OpenMP (planned)

PRINCIPAL COMPONENT ANALYSIS (PCA)

- – Dimensionality reduction technique
- – Transforms high-dimensional data into lower-dimensional space
- – Preserves maximum variance in the data

Eigenfaces Approach

- Represents faces as linear combinations of
base faces
- Reduces facial images to principal components
- Maintains essential features while reducing
data size

Working Process

- Compute mean face from training
images
- Calculate covariance matrix
- Extract eigenvectors (eigenfaces)
- Project faces onto eigenface space

LINEAR DISCRIMINANT ANALYSIS (LDA)

- Supervised learning algorithm ,Focus on maximizing class separability ,Finds optimal projection for classification

Fisherfaces Approach

- Maximizes between-class scatter
- Minimizes within-class scatter
- Better handling of lighting and expression variations

Working Process

- Calculate within-class and between-class scatter matrices
- Compute eigenvectors of scatter matrices
 - Project data onto discriminant space
 - Optimize class separation

APPLICATION IN FACIAL RECOGNITION – PCA – 1

- Converts high-dimensional data into a lower-dimensional space while retaining maximum variance.

1. Data Preprocessing

- Load PGM images and normalize pixel values.
- Compute the mean face vector.
- Center the data by subtracting the mean face vector.

2. Covariance Matrix Computation:

- Calculate the covariance matrix using the centered data.

3. Eigenvector and Eigenvalue Computation

- Use power iteration to compute eigenvalues and eigenvectors.
- Select top 'k' eigenvectors as eigenfaces.

APPLICATION IN FACIAL RECOGNITION – PCA – 2


4. Projection and Recognition

- Project training and test data onto the eigenfaces.
- Use Euclidean distance to find the nearest match in the feature space.


5. Output

- Save the mean face and eigenfaces as PGM files.
- Display recognition accuracy.


PCA Total Execution Time
(Sequential Implementation)



PCA Operation Time Breakdown



PCA Bottleneck Operations Growth



AREAS FOR PARALLELIZATION – PCA 1

1. Normalization of Pixel Values

- Normalize images by dividing each pixel by the maximum value.
- Parallelization: Distribute the computation of pixel normalization across threads (OpenMP) or CUDA blocks.

2. Mean Vector Computation

- Compute the mean of all training images pixel-wise.
- Parallelization: Sum the pixel values across images in parallel.

3. Centering the Data

- Subtract the mean vector from each image.
- Parallelization: Process each image independently.

AREAS FOR PARALLELIZATION – PCA 2

4. Covariance Matrix Computation

- Calculate the covariance matrix using matrix multiplication.
- Parallelization: Perform matrix multiplications in parallel using OpenMP or CUDA.


5. Eigenvalue and Eigenvector Computation

- Use power iteration to compute eigenvalues and eigenvectors.
- Parallelization: Parallelize the matrix-vector multiplications within each iteration.


6. Projection onto Eigenfaces

- Project the training and test data onto the top 'k' eigenvectors.
- Parallelization: Process each image independently during projection.

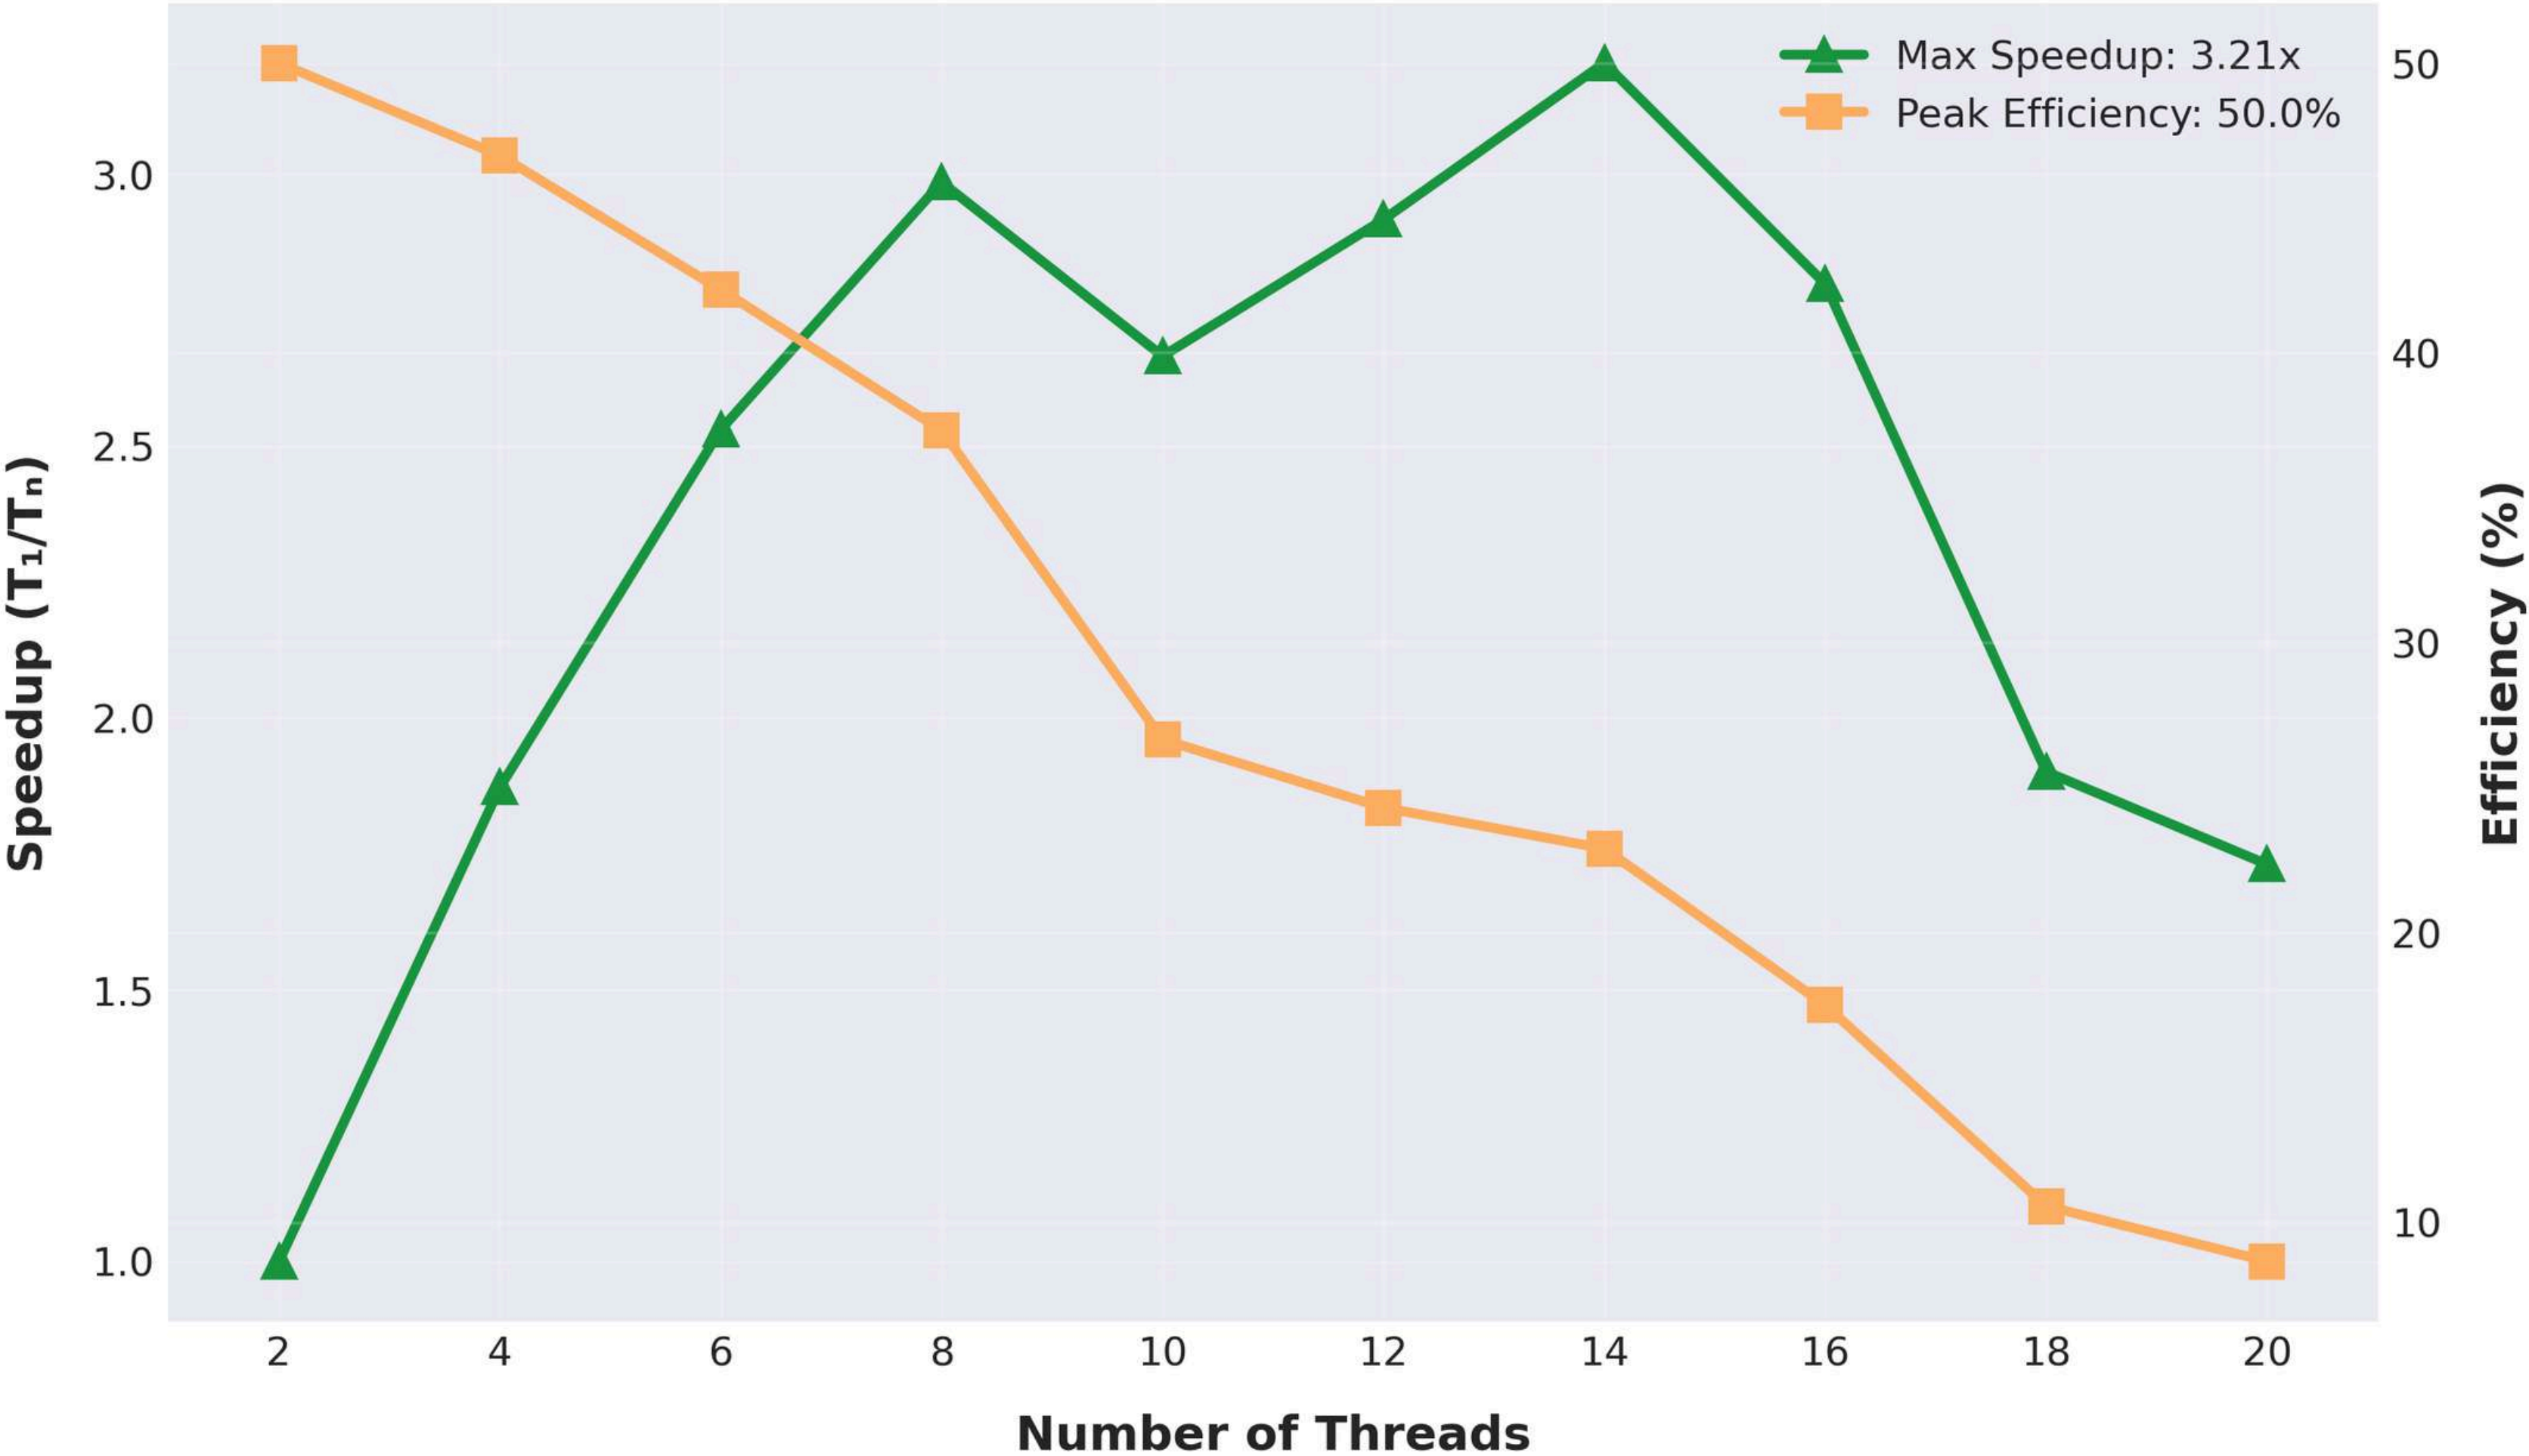
Parallel Scaling Performance



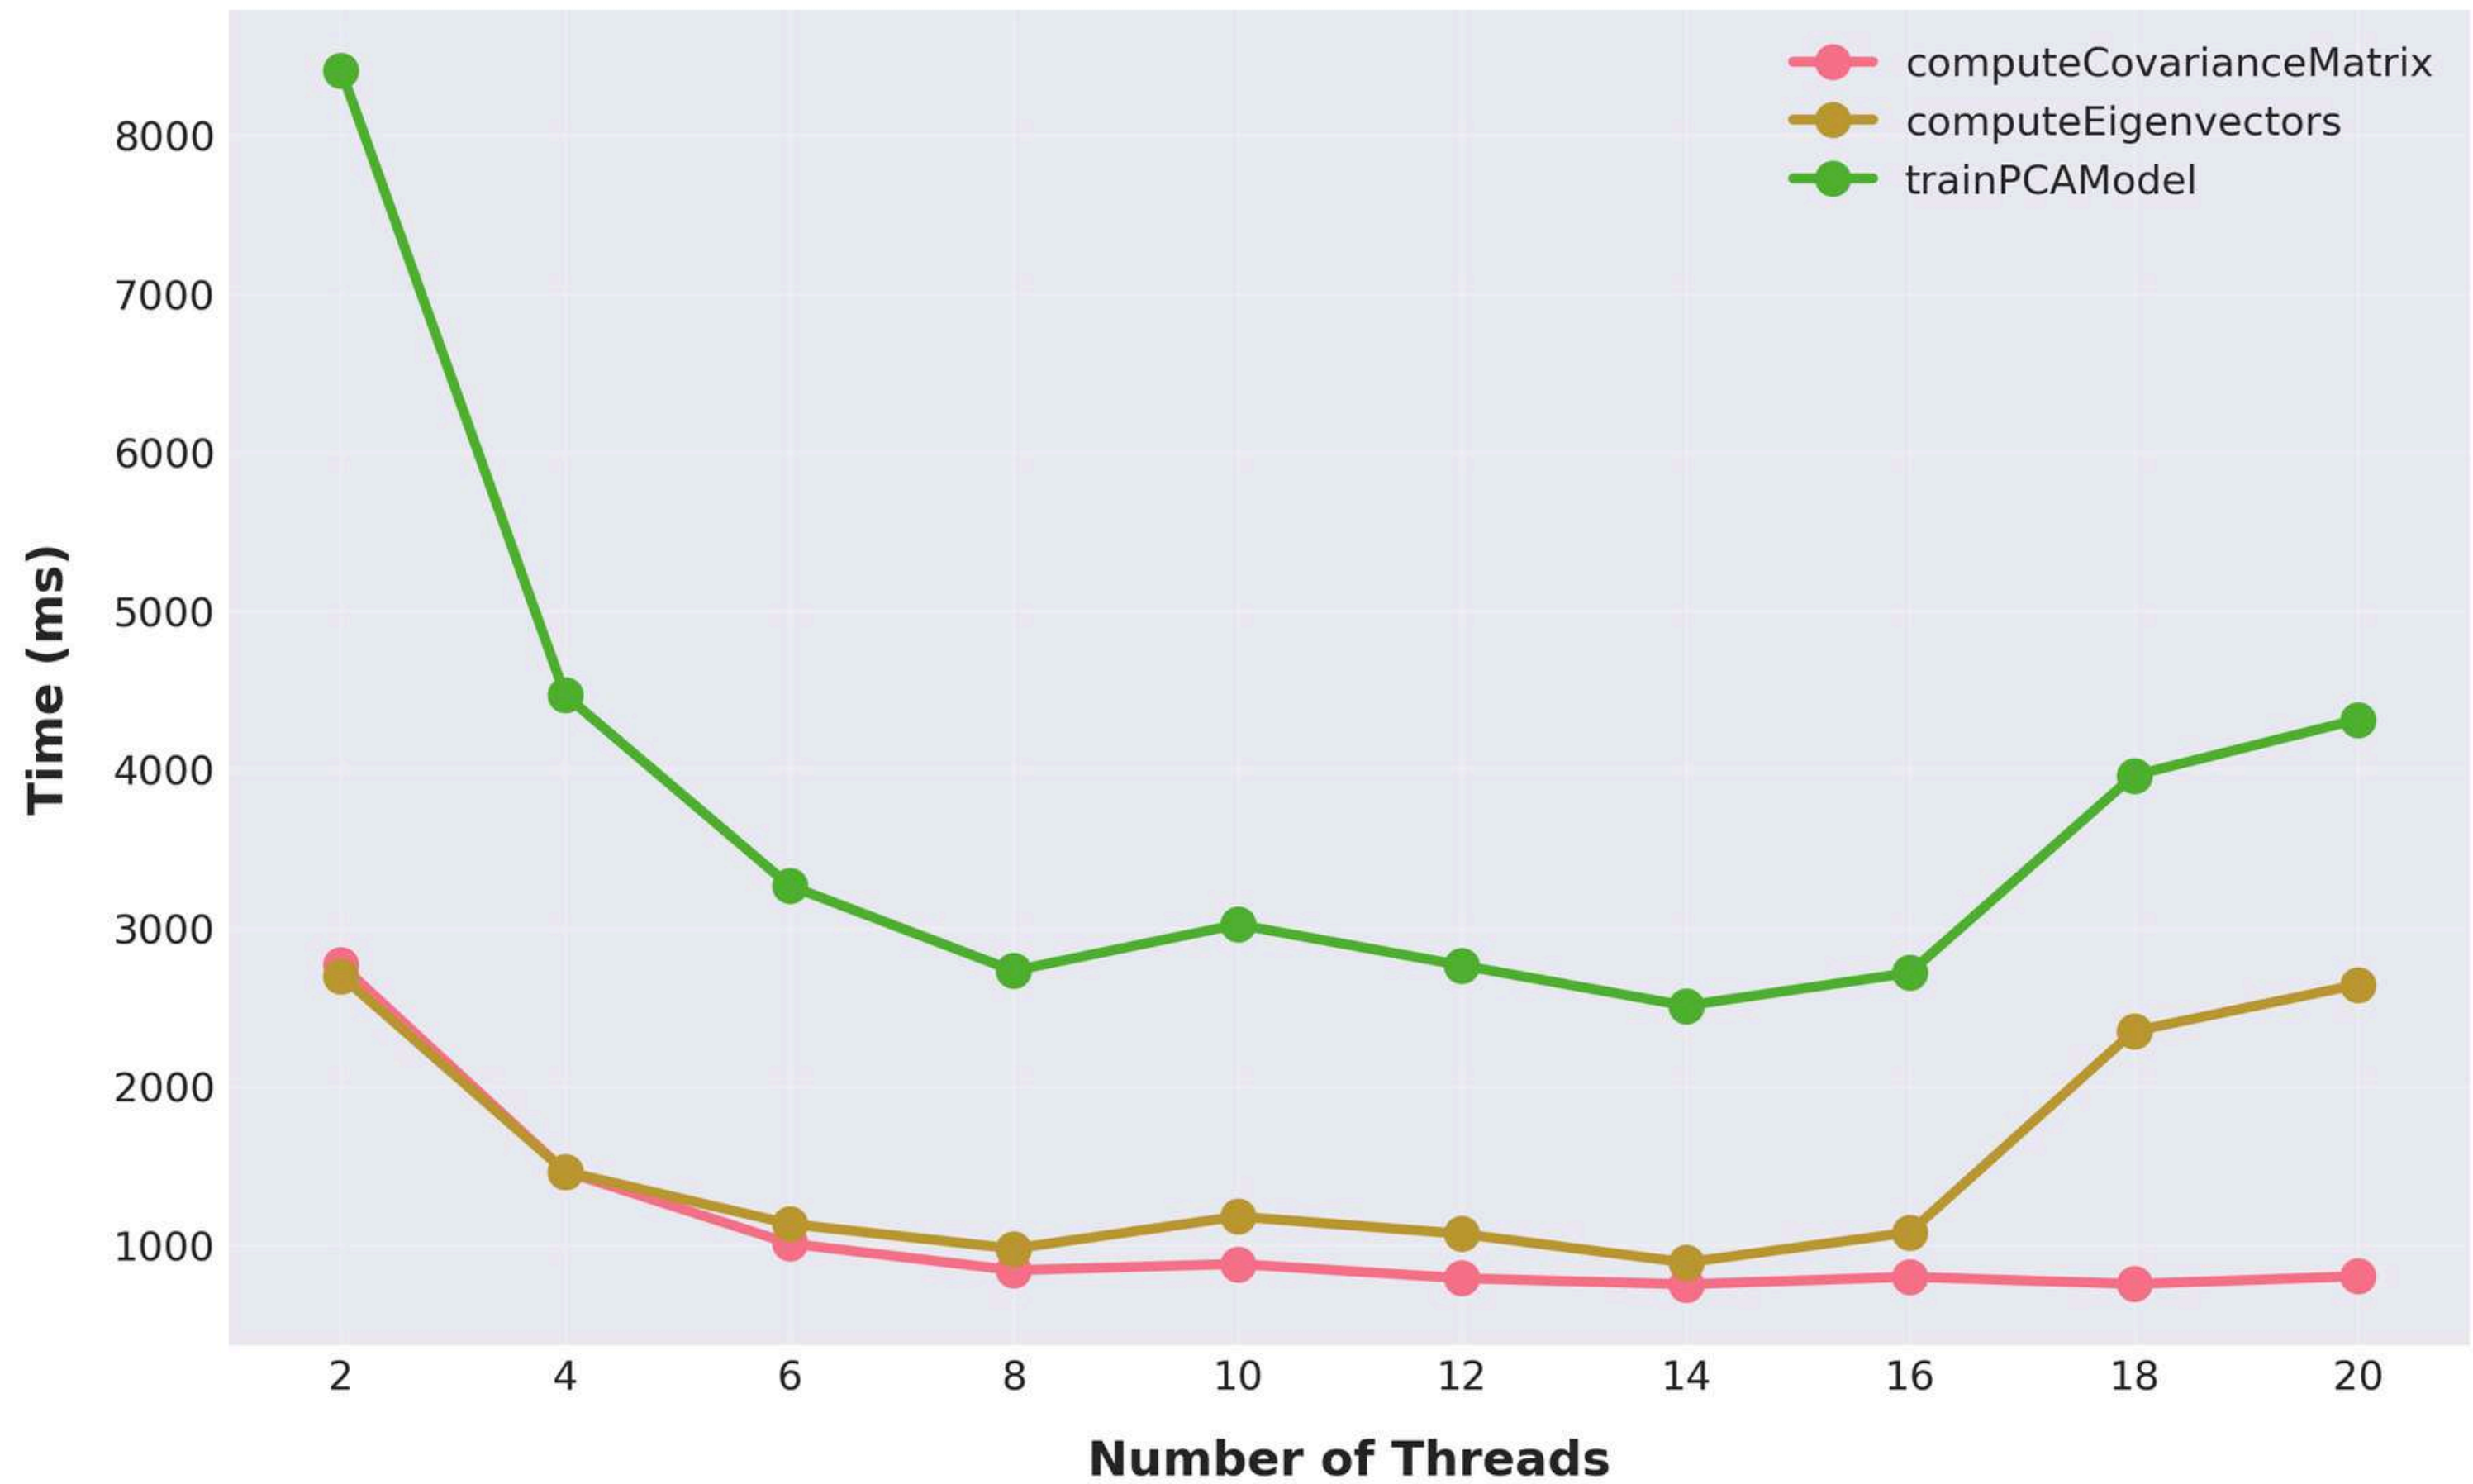
Operation Time Distribution




Parallel Speedup Analysis




Bottleneck Operation Scaling




PCA Total Execution Time
(Parallel Implementation)




PCA Operation Time Breakdown (OMP)



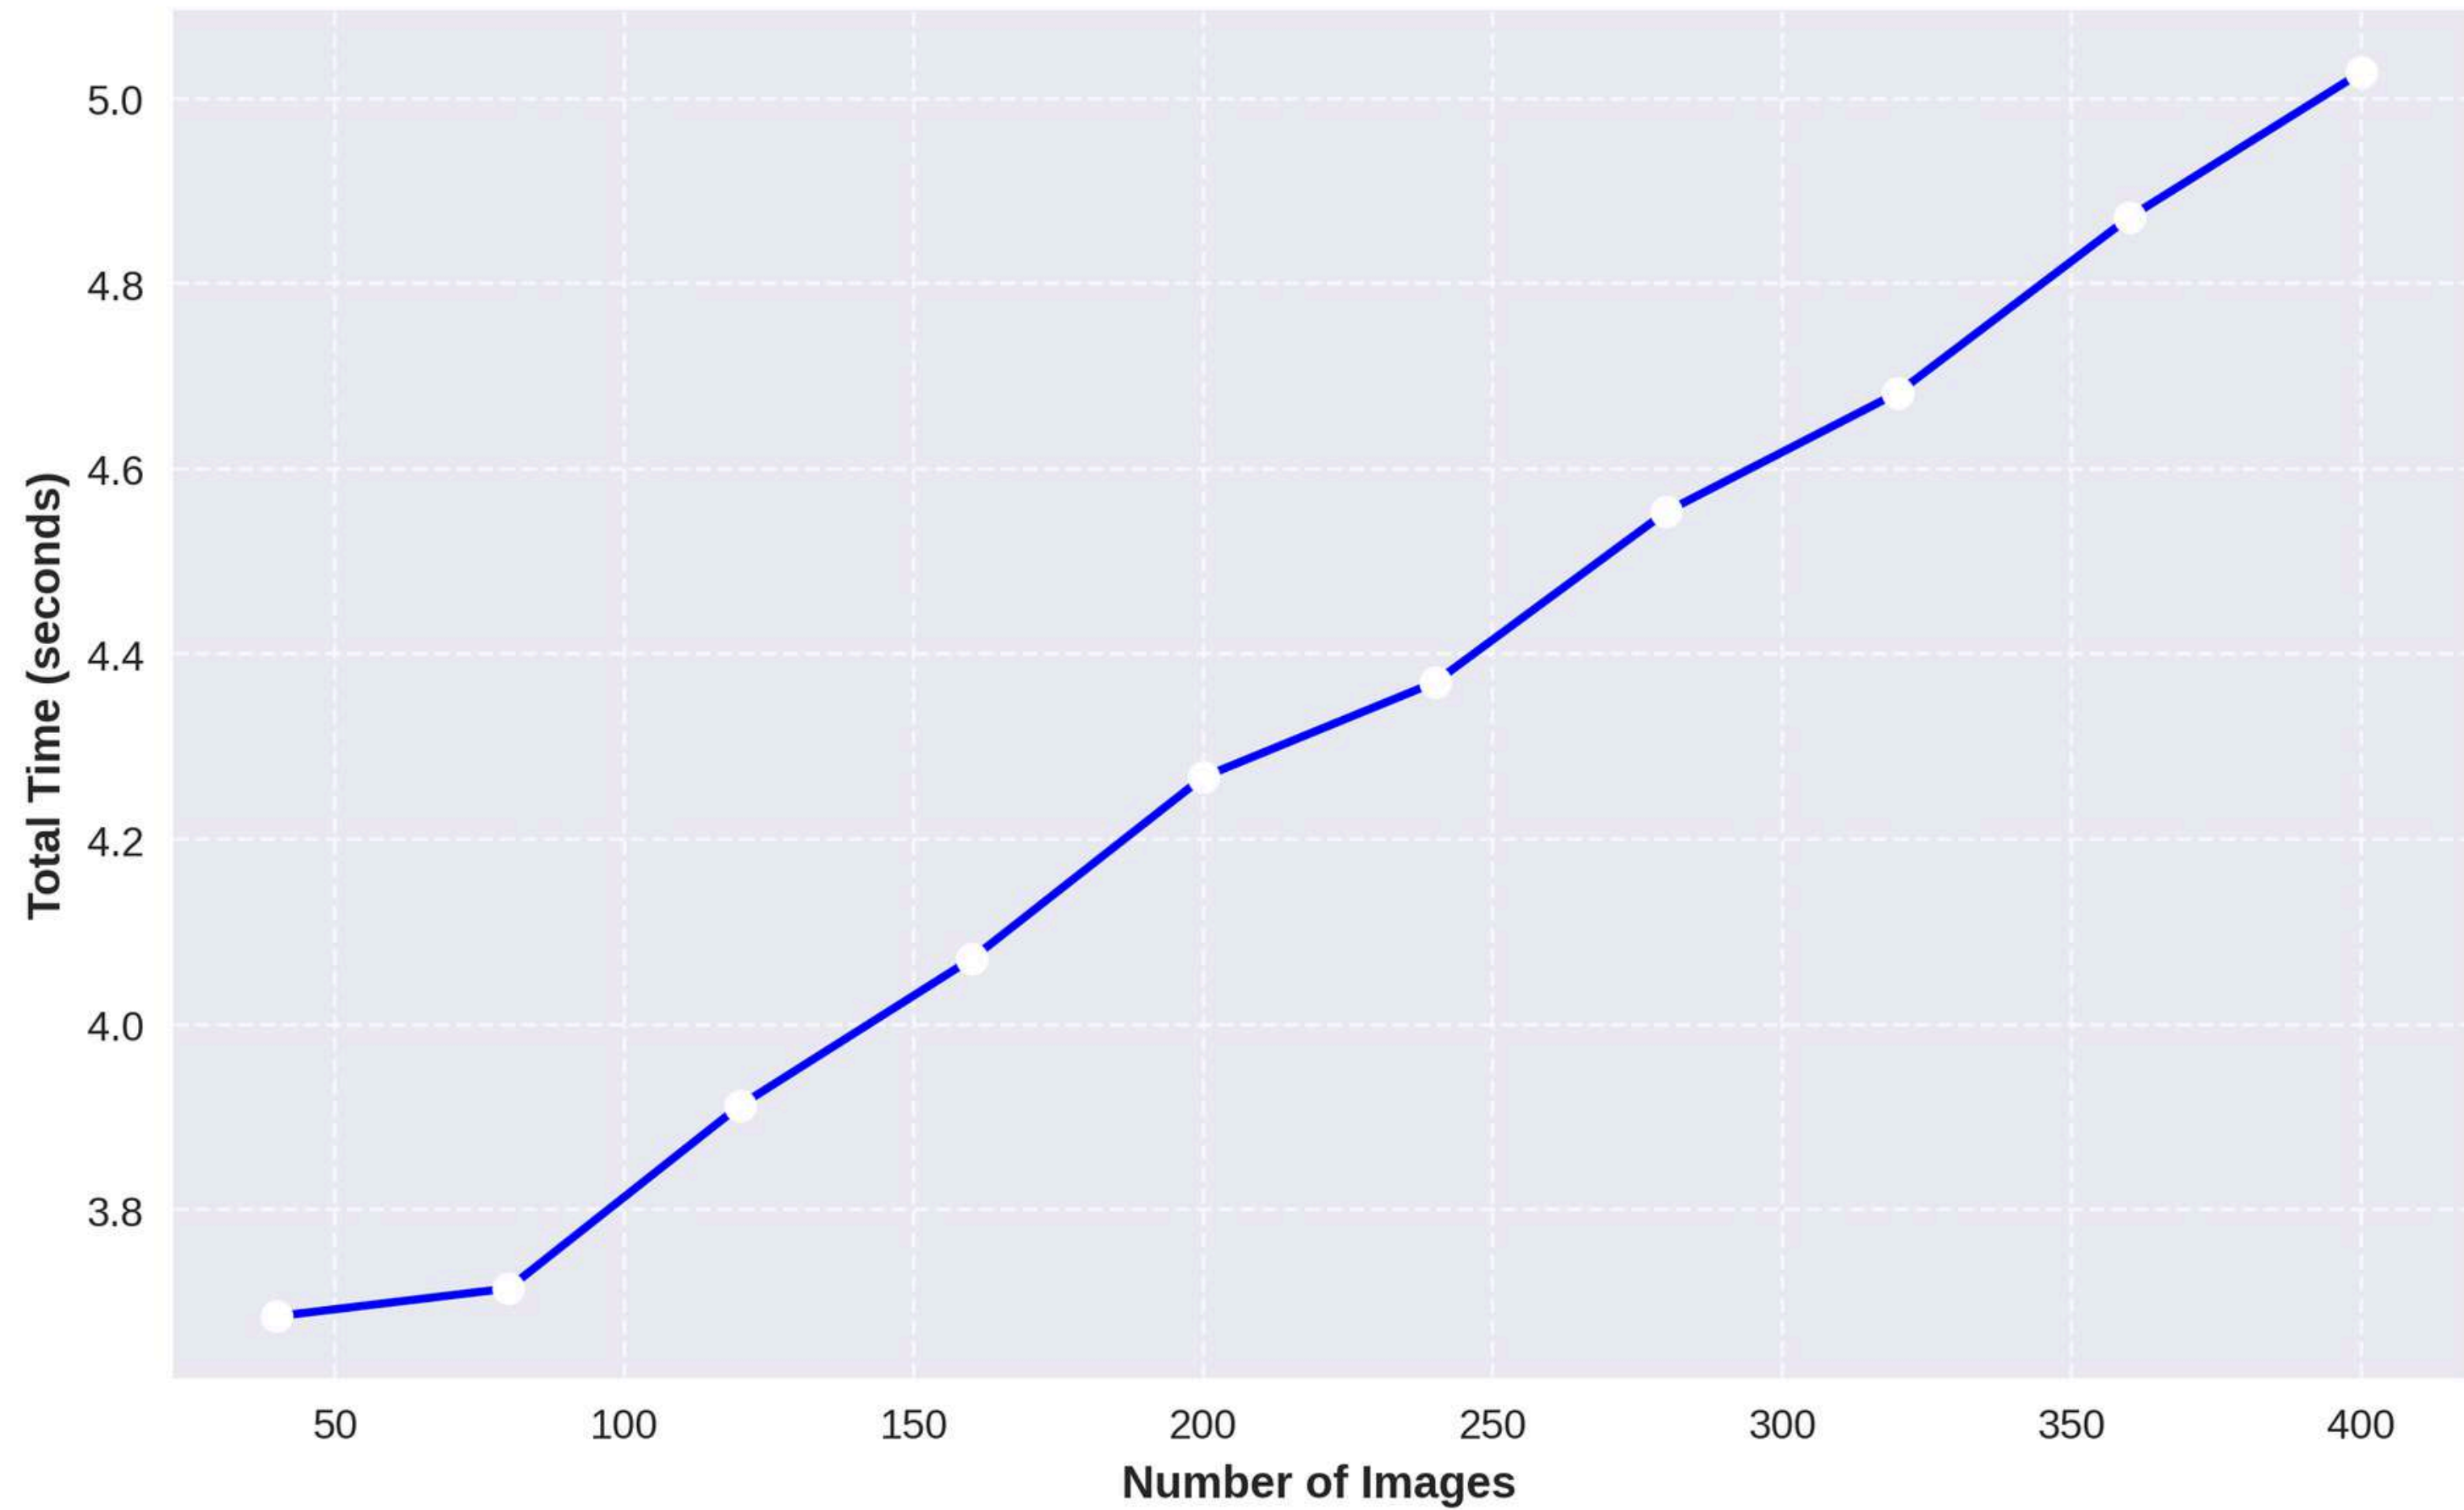
PCA Efficiency: Time per Image (OMP)




PCA Bottleneck Operations Growth (OMP)



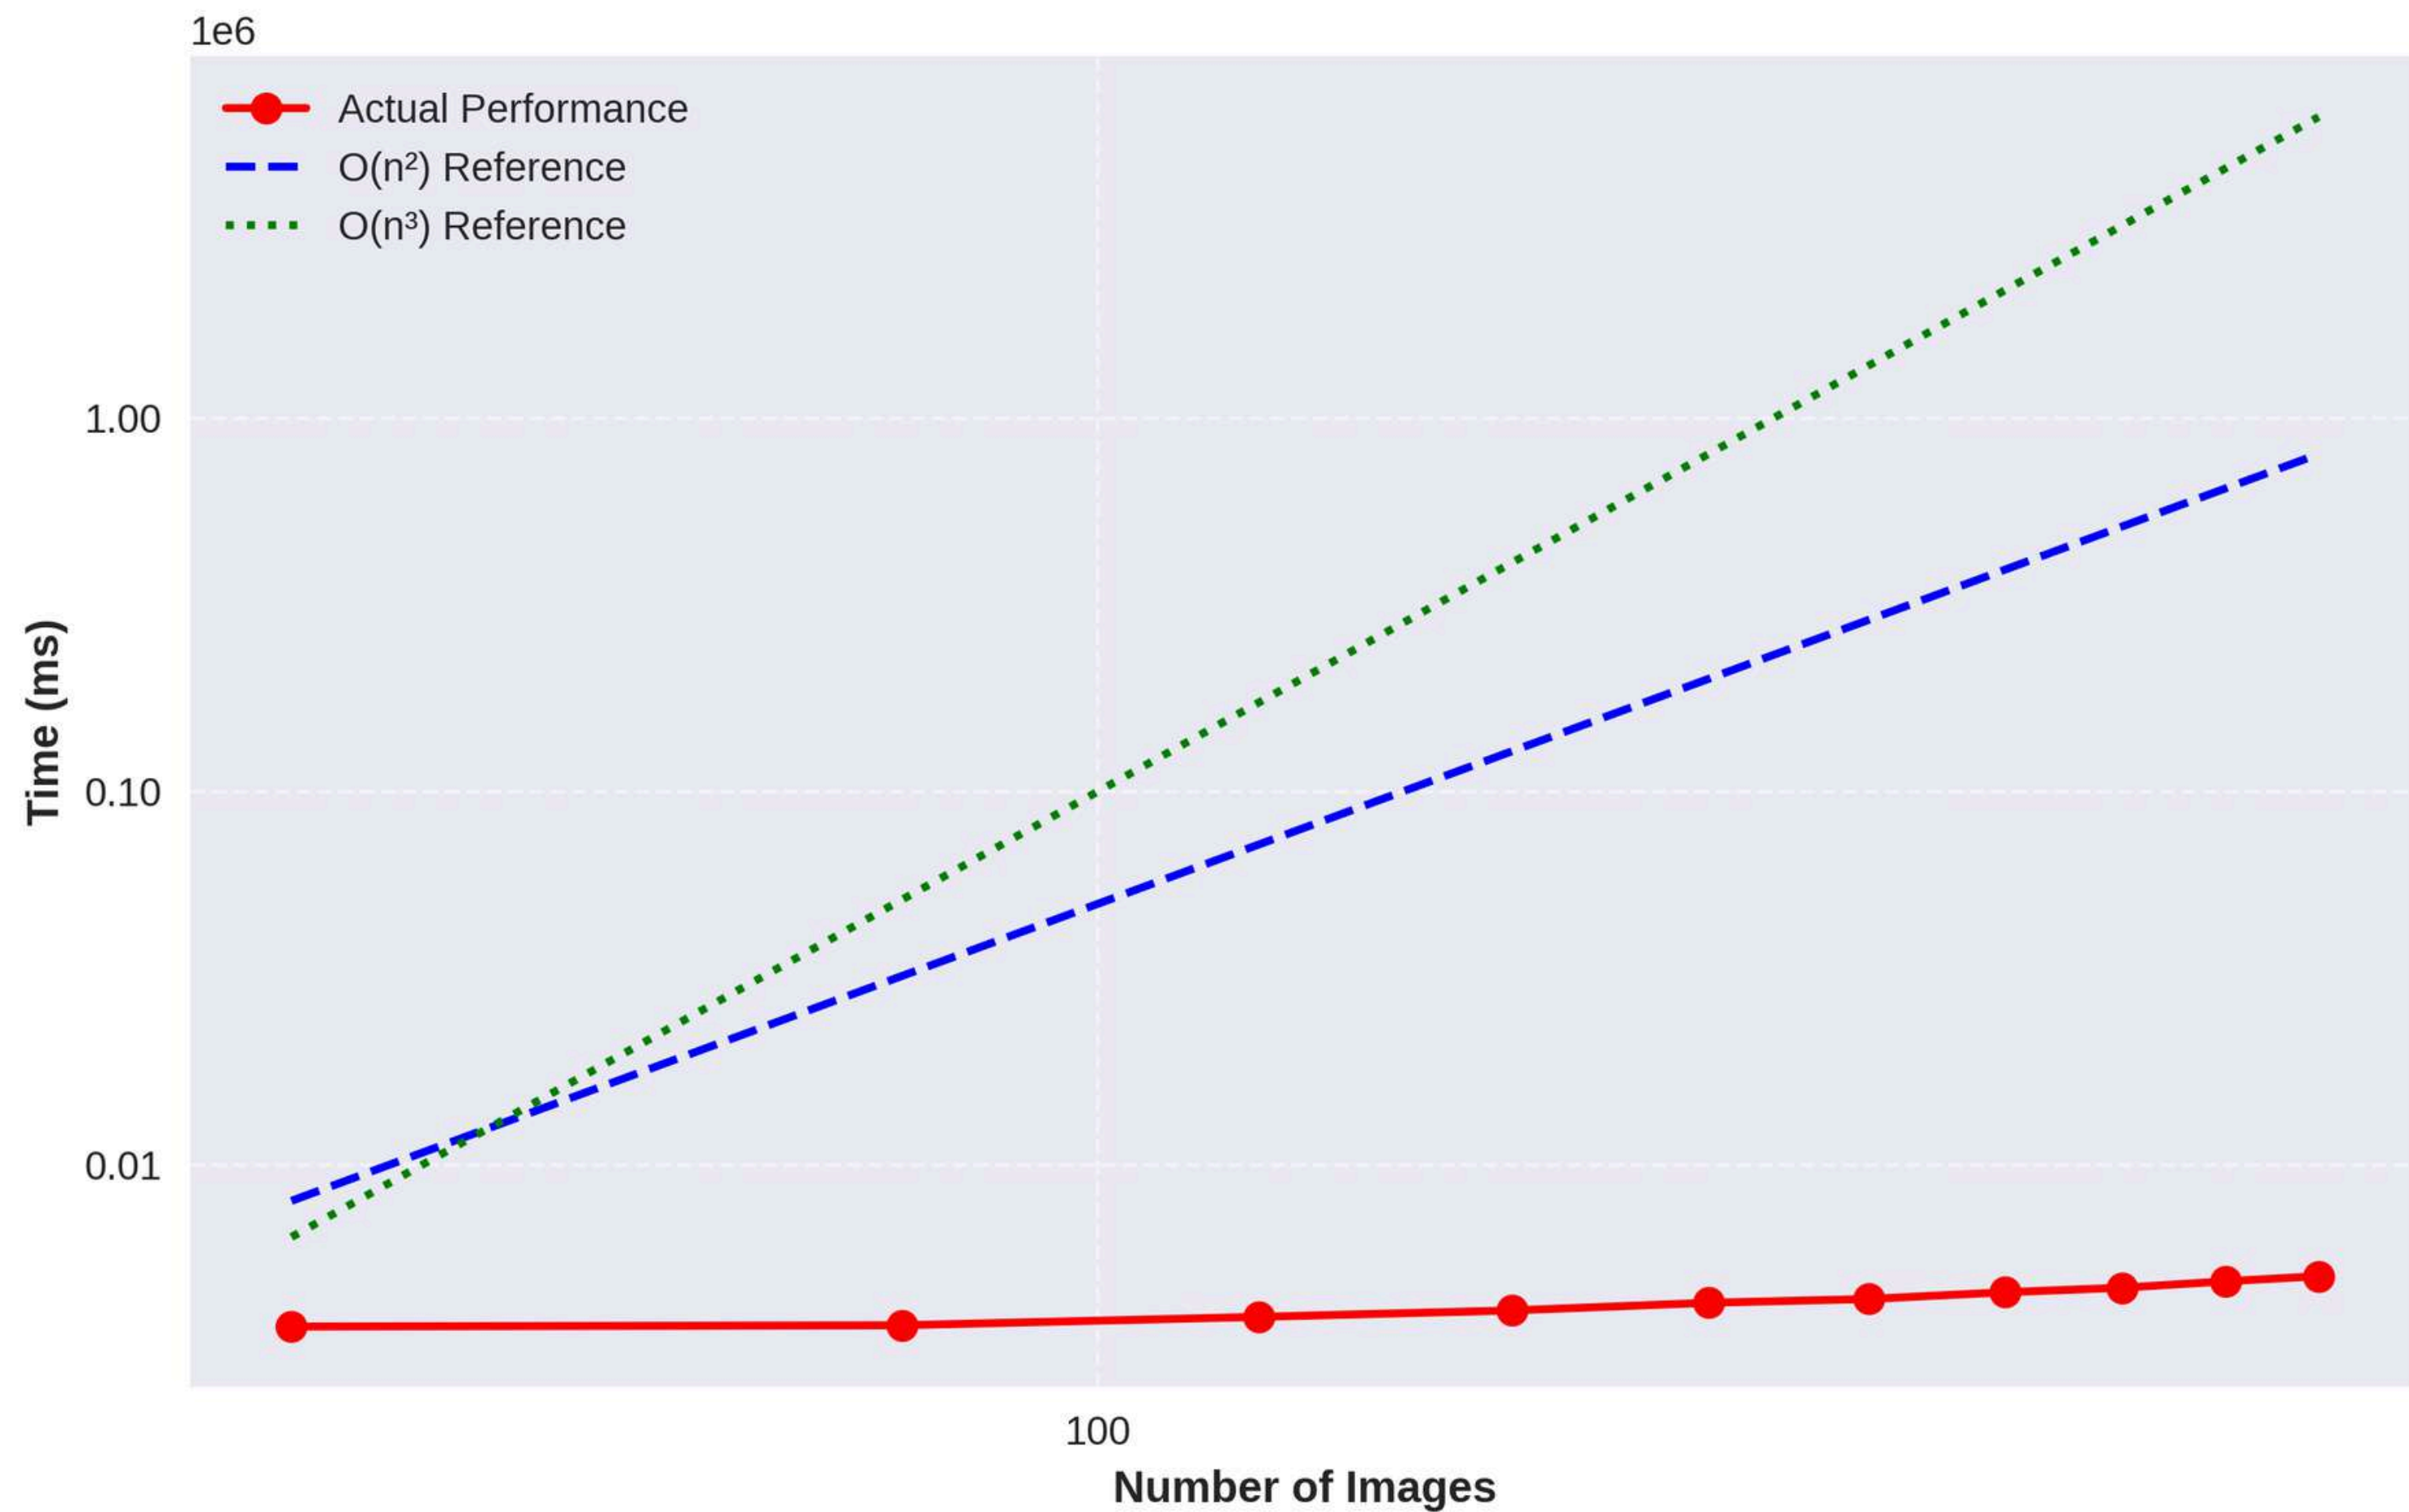
PCA Total Execution Time
(CUDA Implementation)




PCA Efficiency: Time per Image (CUDA)




PCA Time Complexity Analysis (CUDA)



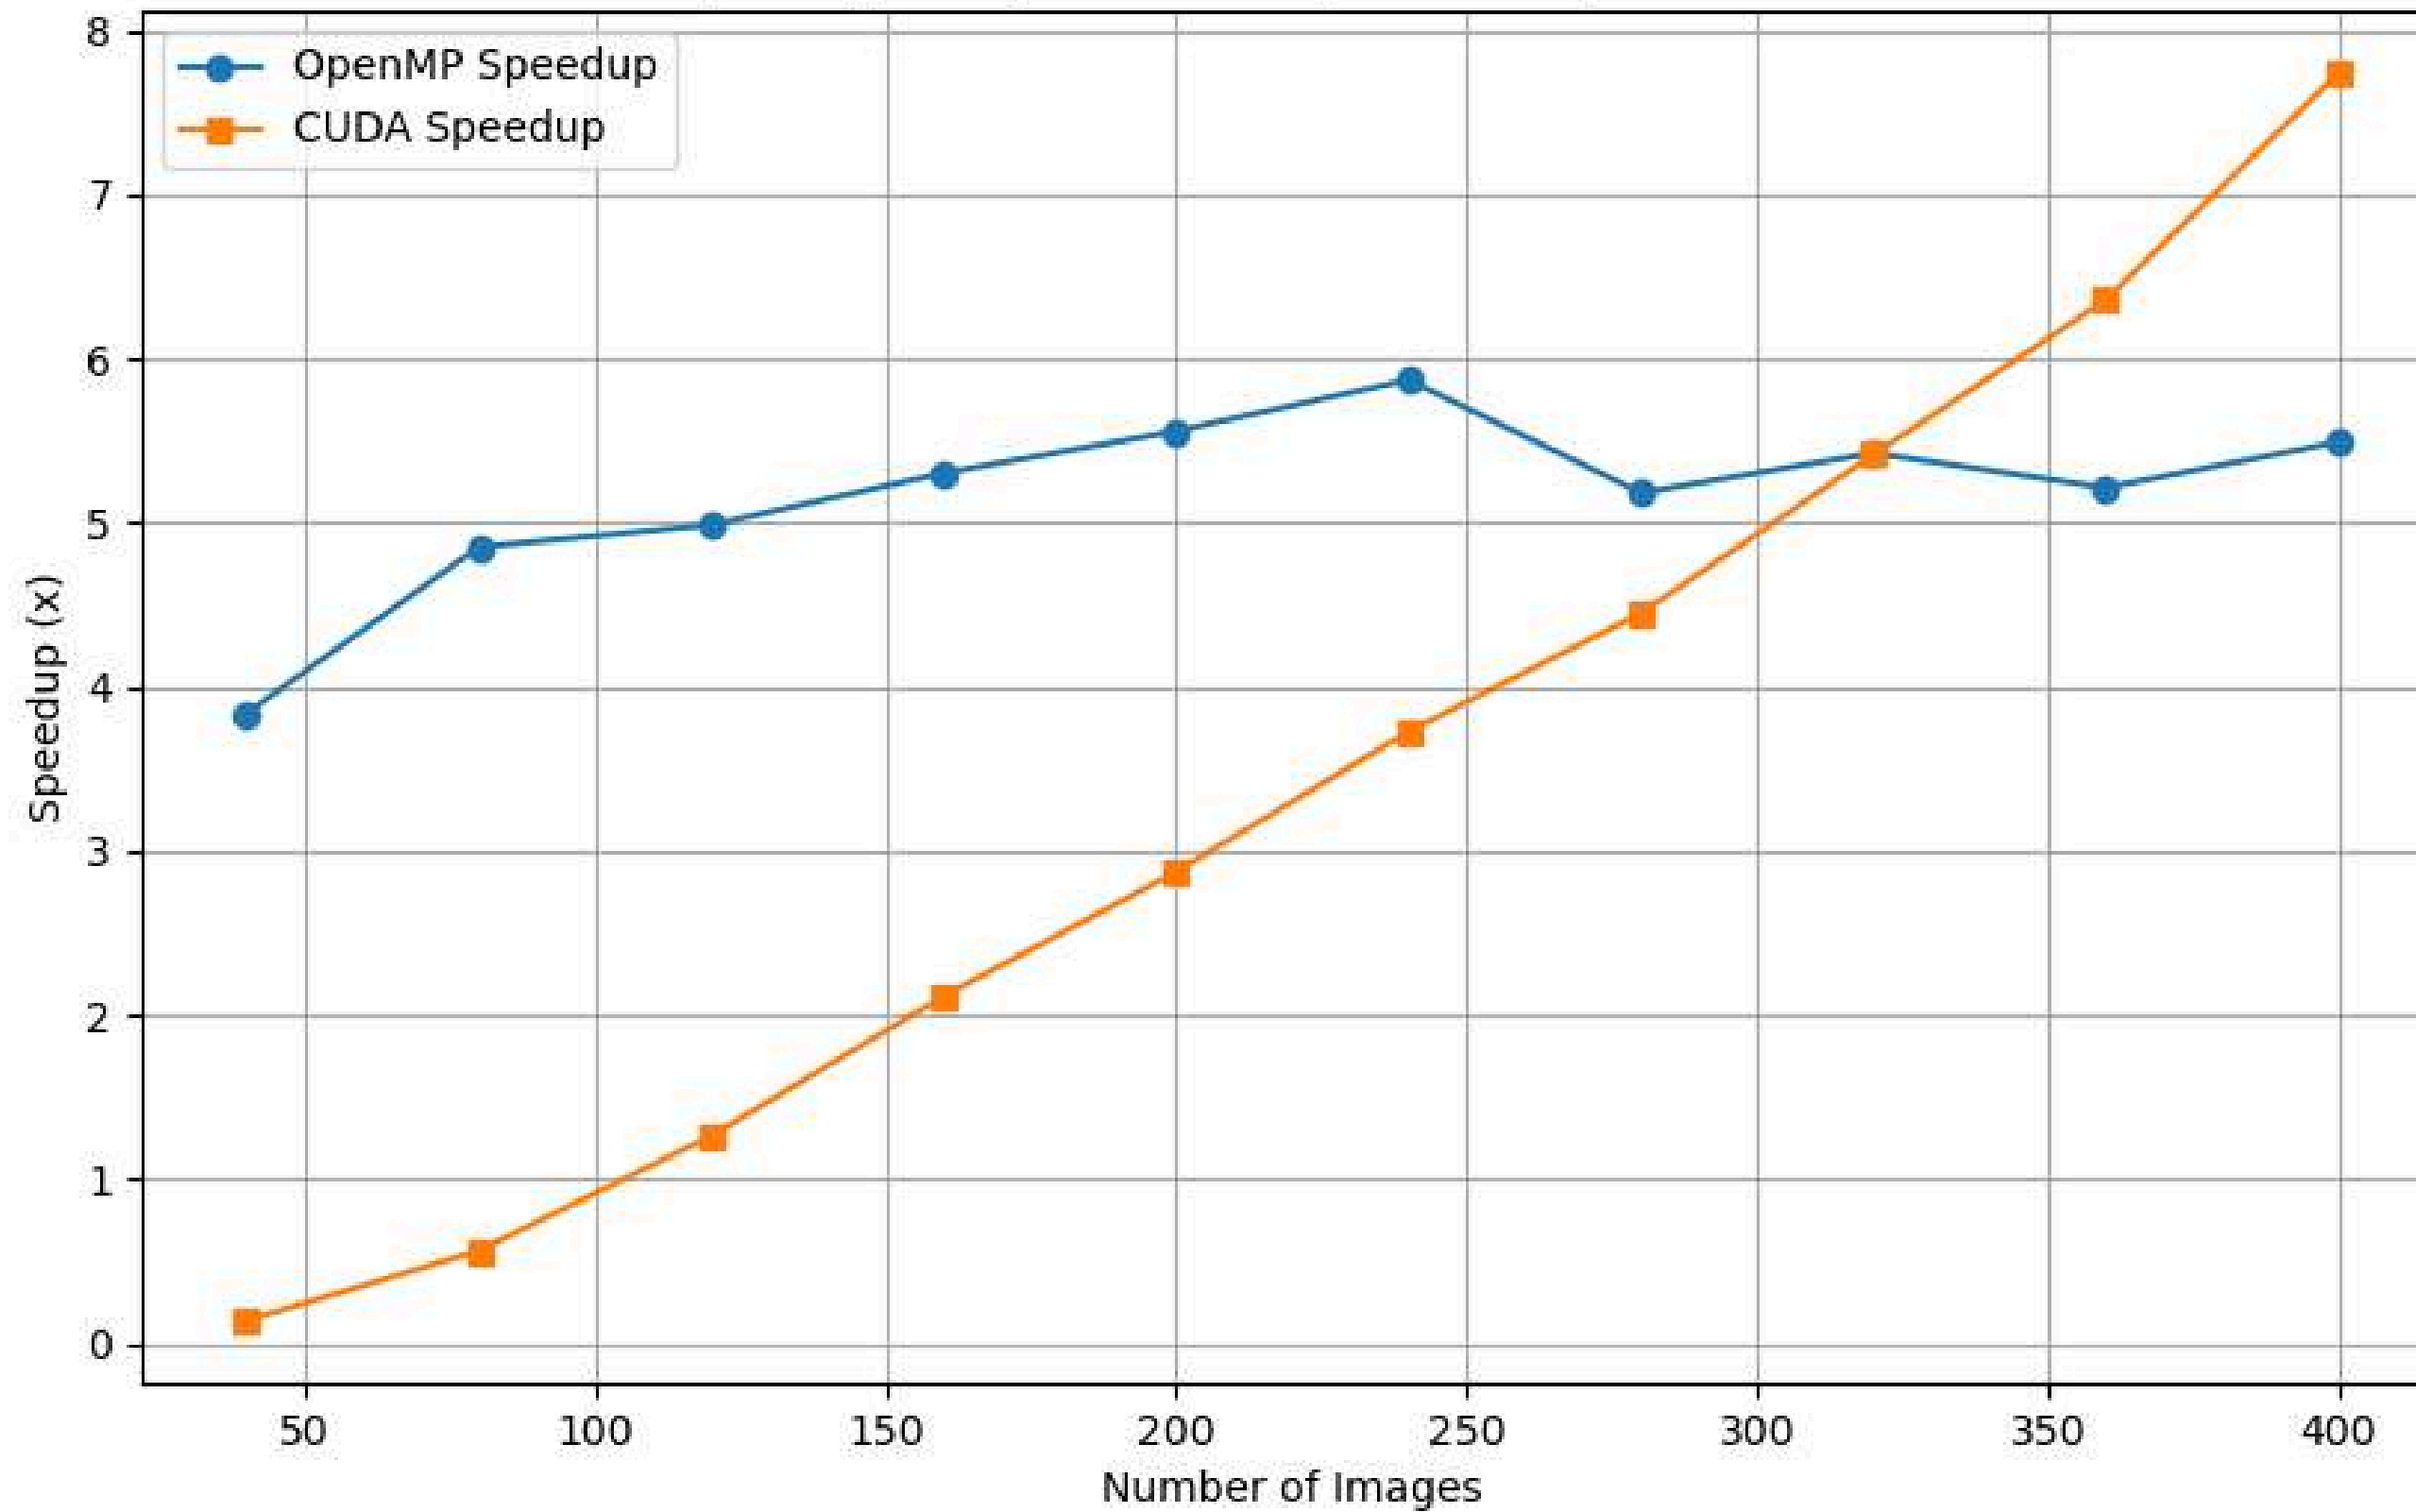
PCA Bottleneck Operations Growth (CUDA)




Total Execution Time Comparison



Speedup Compared to Sequential Implementation



Time Breakdown Comparison for 400 Images



APPLICATION IN FACIAL RECOGNITION – LDA – 1

- Supervised dimensionality reduction technique.
- Maximizes class separability while minimizing within-class variance.

1. Data Preprocessing

- Load PGM images and normalize pixel values.
- Compute the mean face vector.
- Center the data by subtracting the mean face vector.

2. Class-Specific Data Handling

- Group samples by class (person ID).
- Compute within-class and between-class scatter matrices.

3. PCA as Preprocessing

- Apply PCA to reduce dimensionality before applying LDA.

4. LDA Projection Matrix

- Solve the generalized eigenvalue problem for S_B and S_W .
- Compute Fisherfaces in the PCA-reduced space.

5. Projection and Recognition

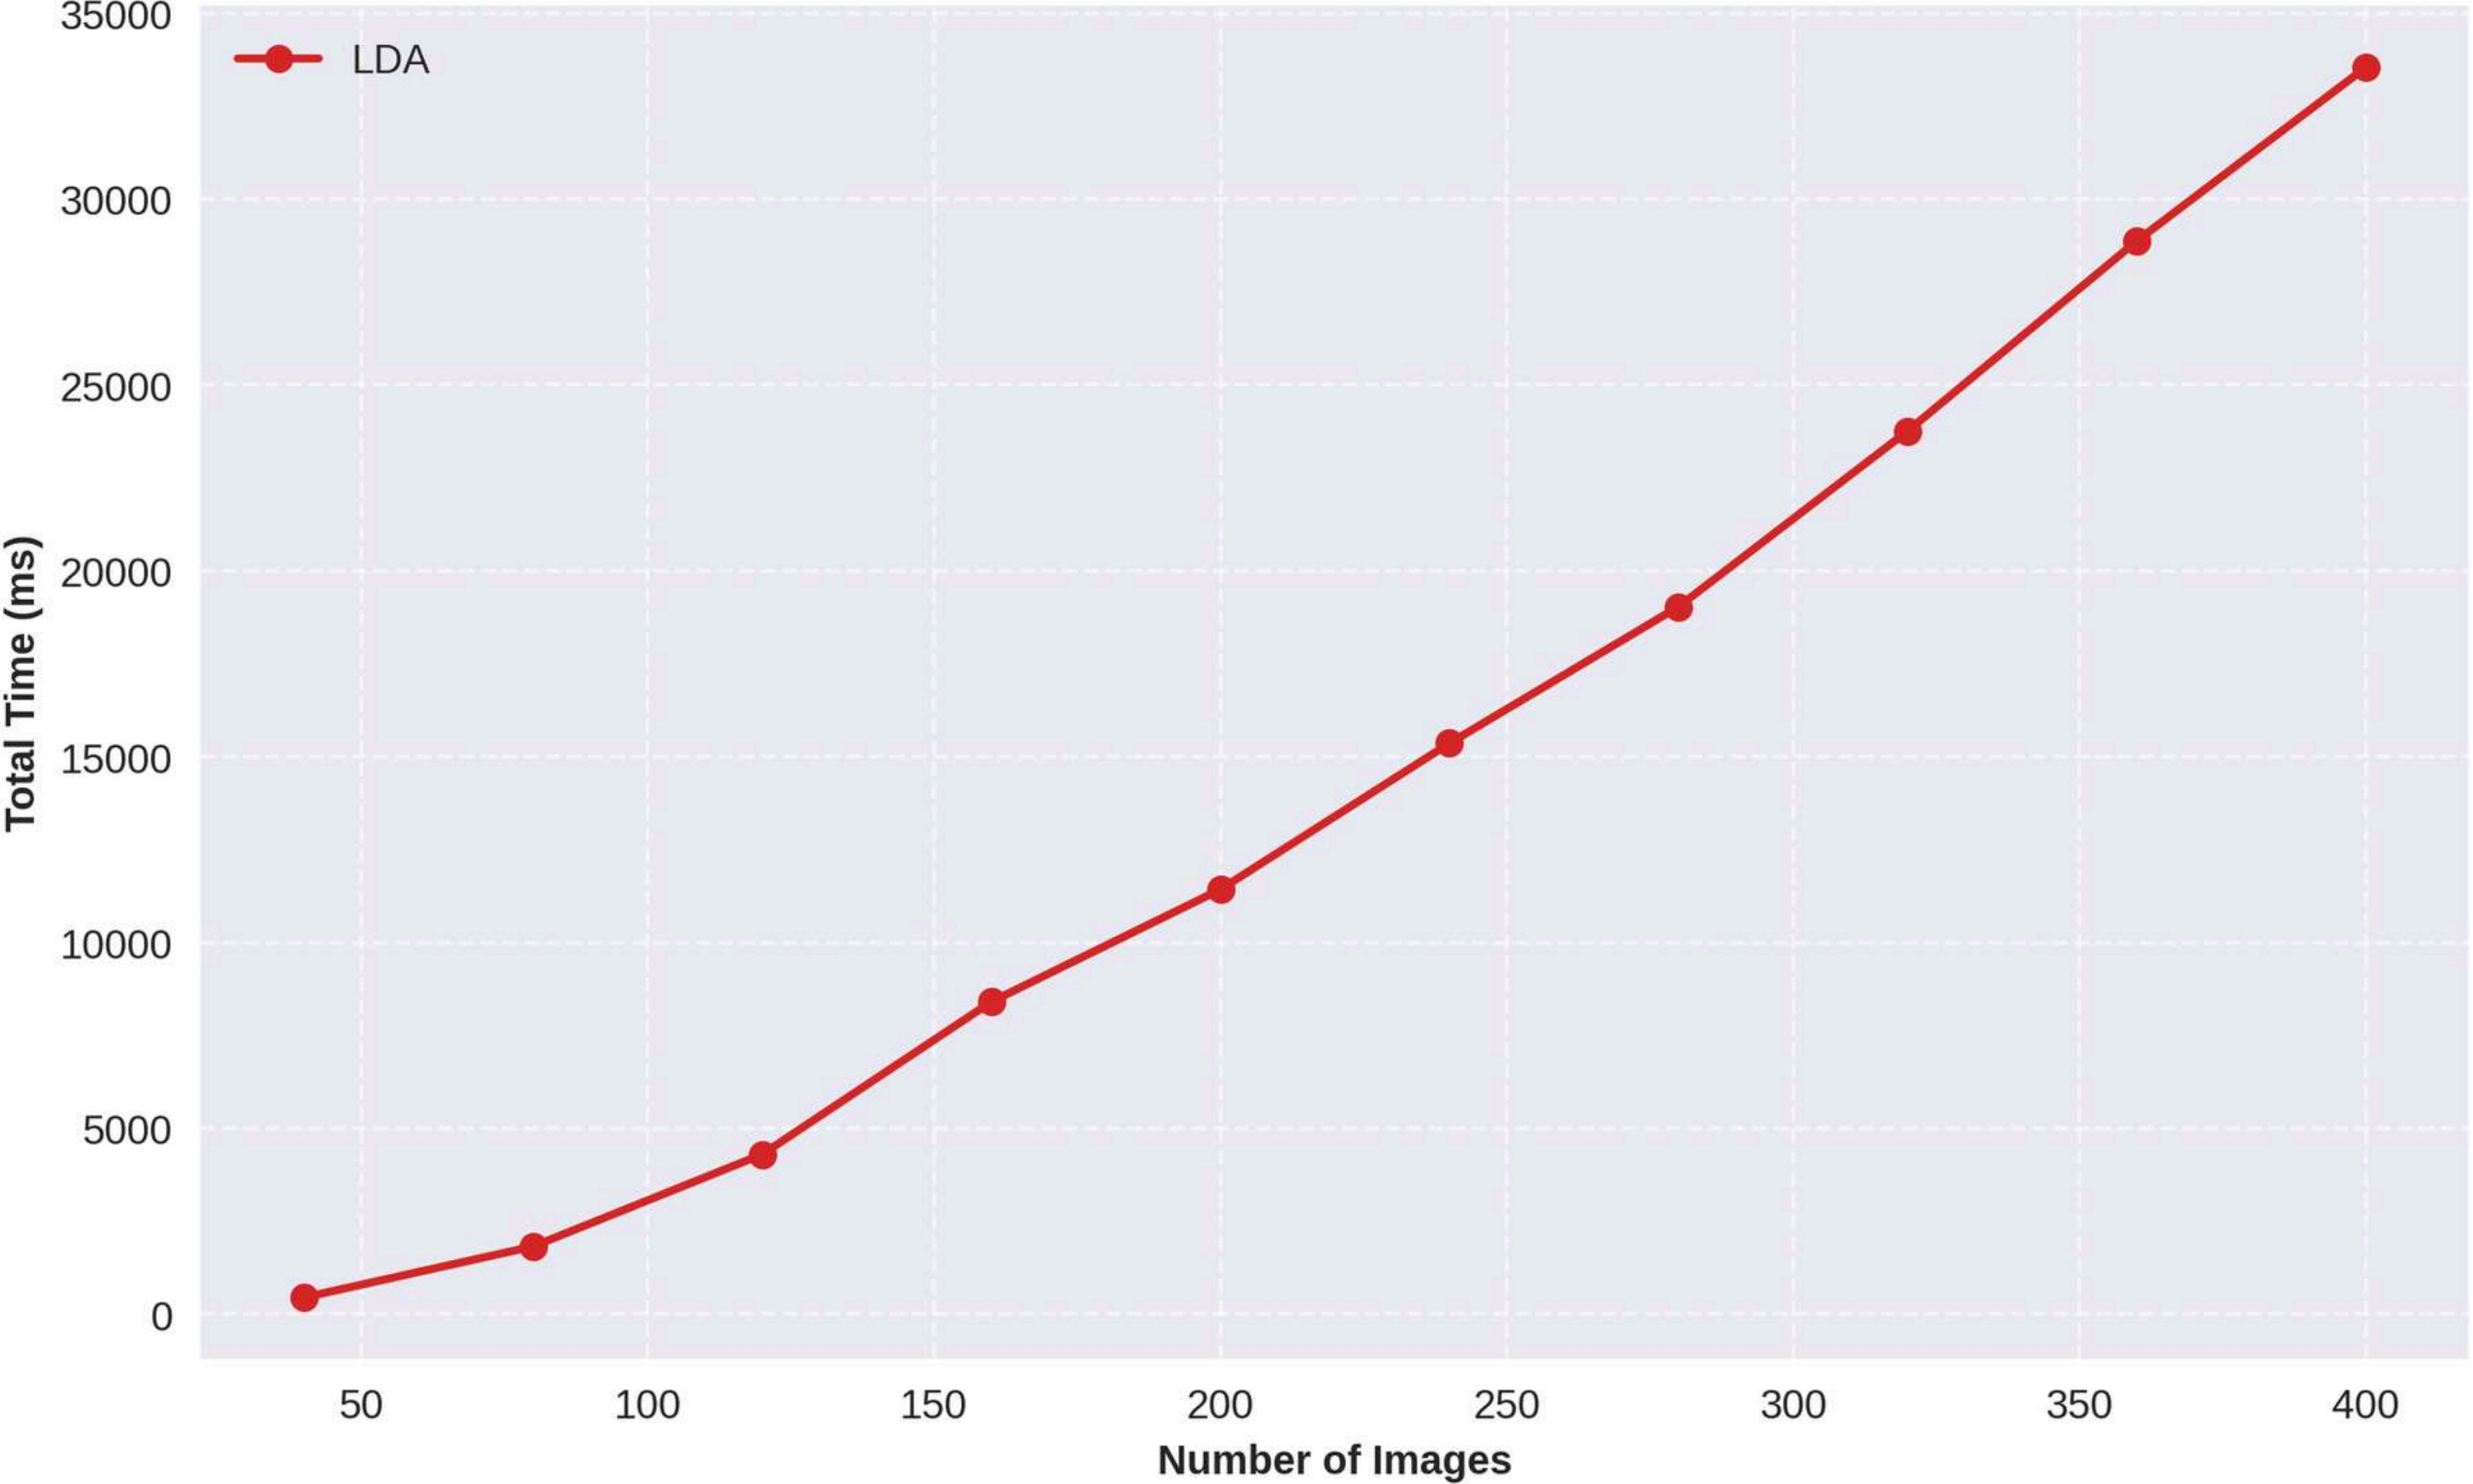
- Project training and test data onto Fisherfaces.
- Use Euclidean distance to find the nearest match.

6. Output


- Save the global mean face and Fisherfaces as PGM files.
- Display recognition accuracy and confusion matrix.

APPLICATION IN FACIAL RECOGNITION – LDA – 2

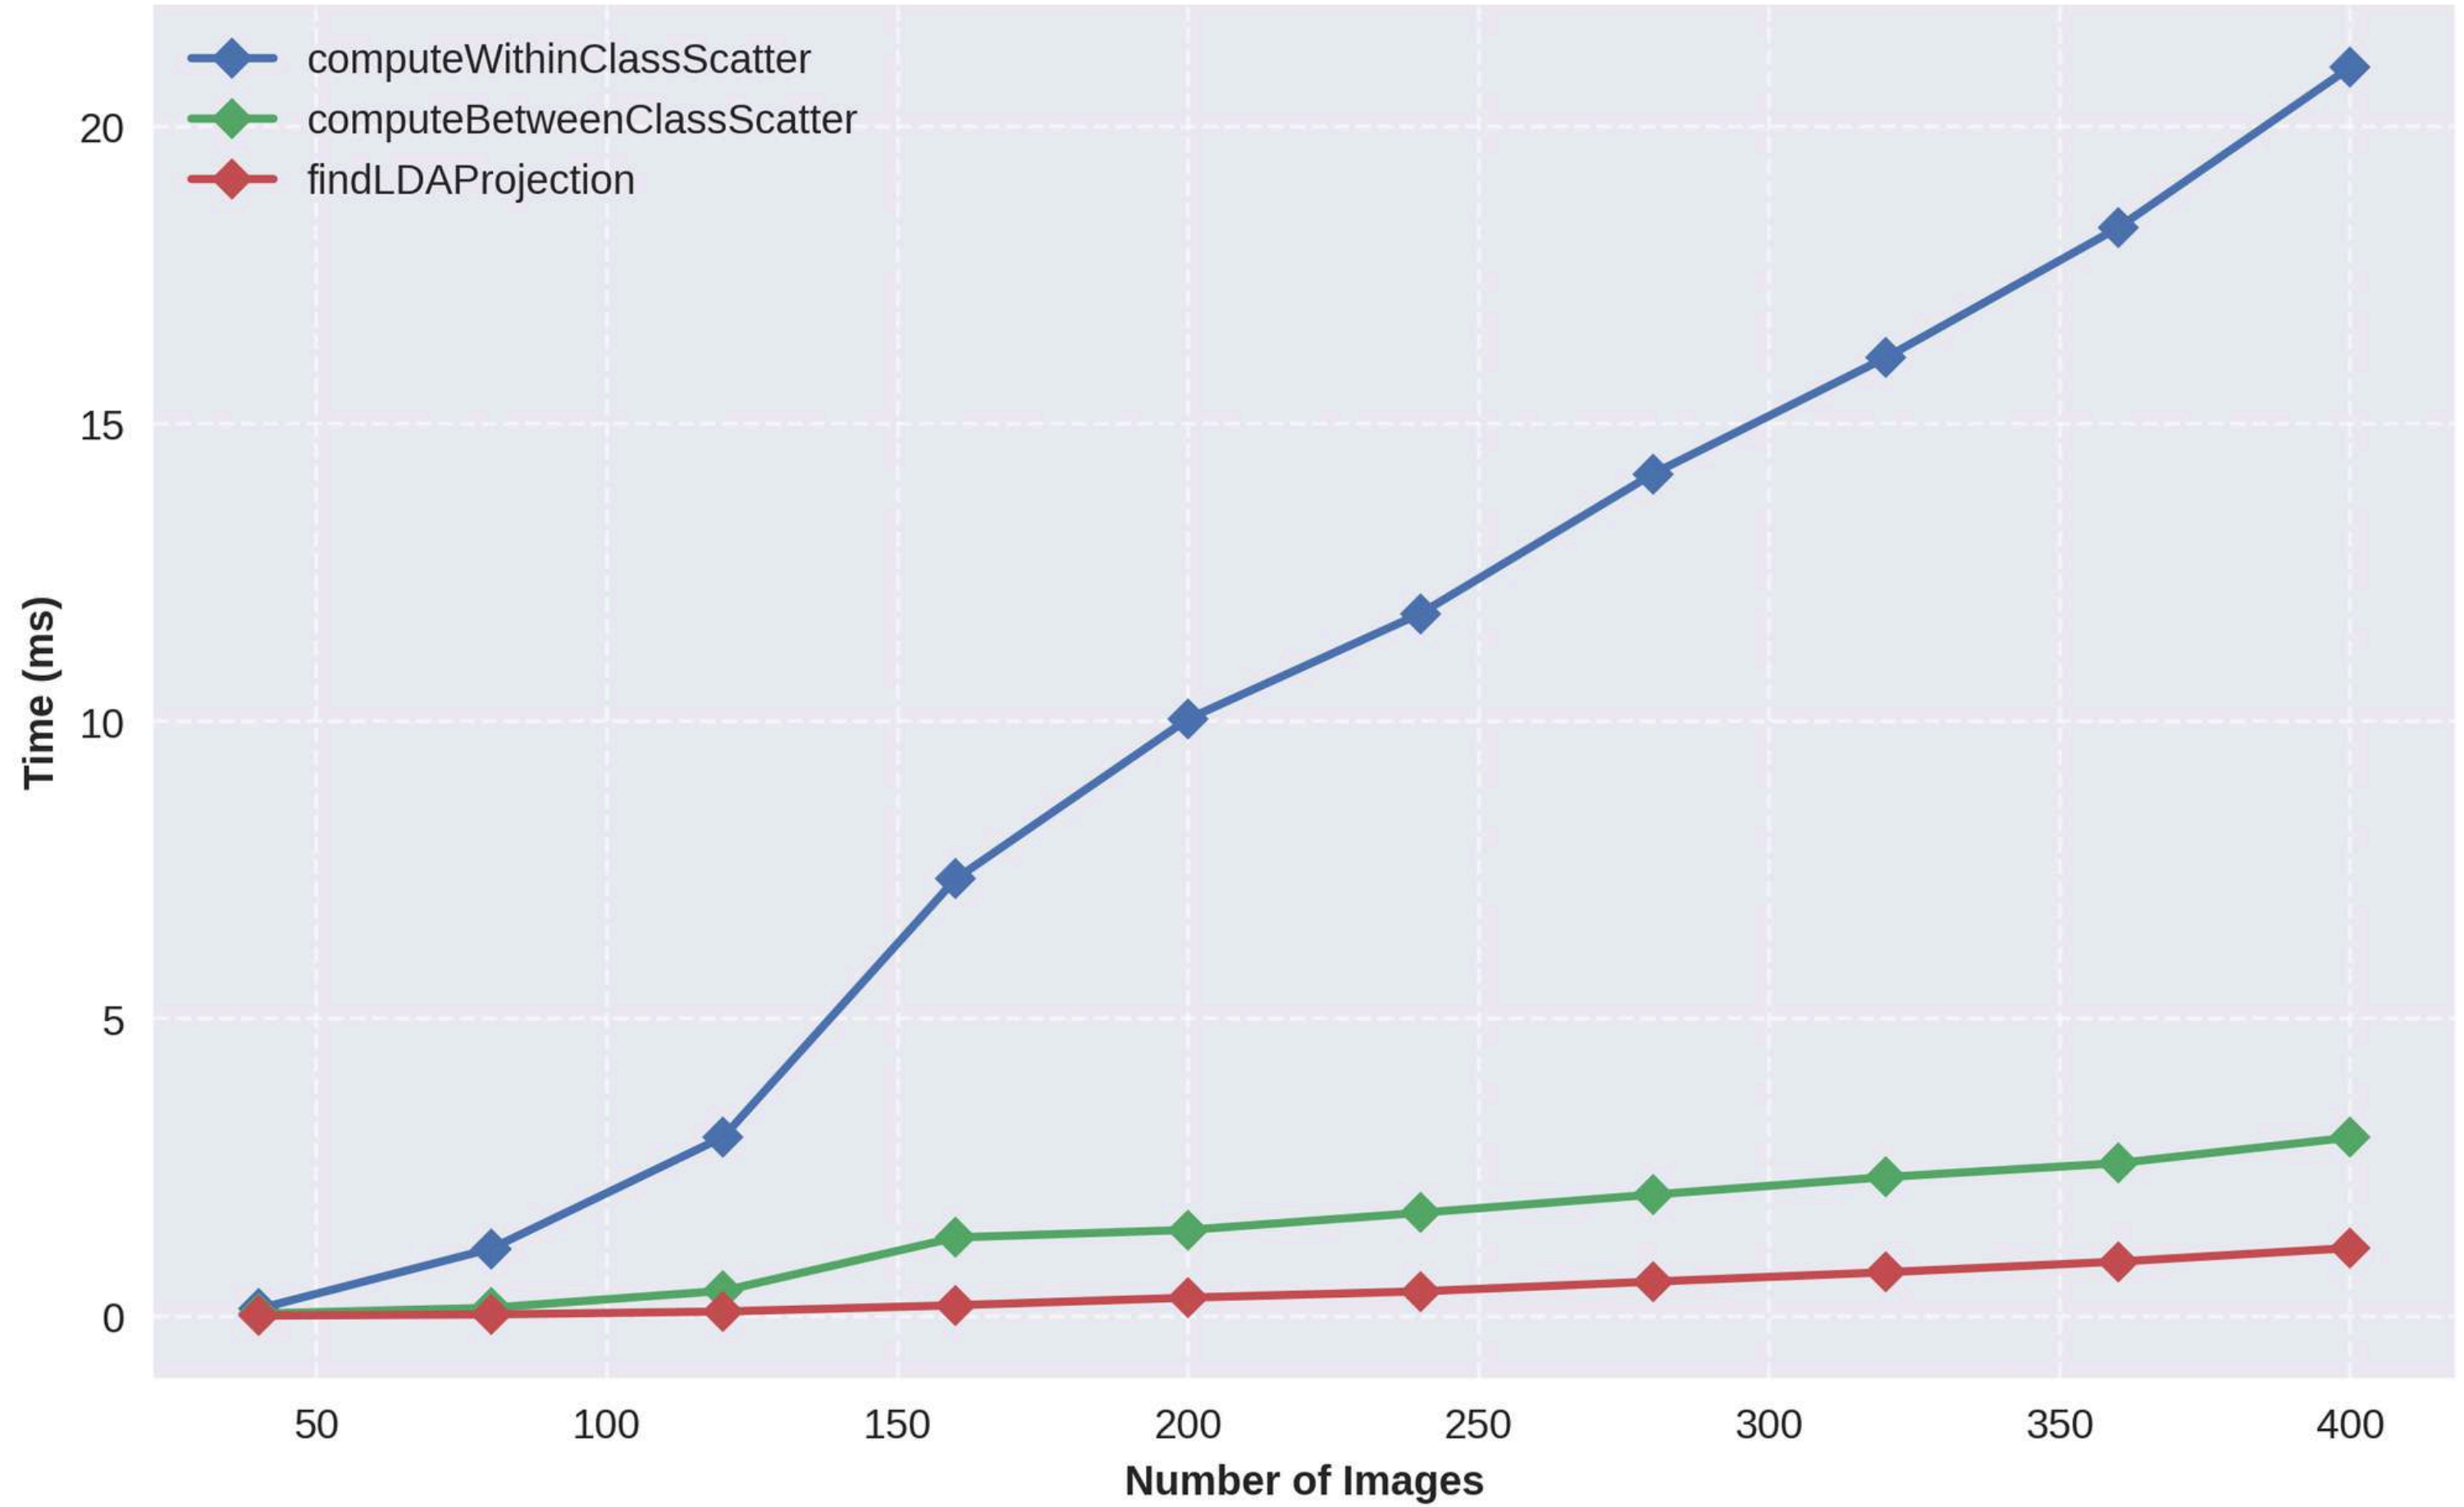
Total Execution Time for LDA Implementation



LDA Performance vs Memory Usage



LDA-Specific Operation Times



AREAS FOR PARALLELIZATION – LDA 1

1. Normalization of Pixel Values

- Normalize images by dividing each pixel by the maximum value.
- Parallelization: Similar to PCA, distribute computation across threads or CUDA blocks.

2. Global Mean Computation

- Compute the global mean vector for all training images.
- Parallelization: Sum the pixel values across images in parallel.

3. Within-Class Scatter Matrix (SW) Computation

- Compute scatter matrices for each class and sum them.
- Parallelization: Process each class independently, and parallelize within-class computations.

AREAS FOR PARALLELIZATION – LDA 2

4. Between-Class Scatter Matrix (SB) Computation

- Compute the scatter matrix for the difference between class means and the global mean.
- Parallelization: Each class's computation can run independently.

5. PCA Preprocessing

- Reduce dimensionality using PCA before LDA.
- Parallelization: Apply the same parallelization techniques as in PCA.

6. LDA Projection Matrix Computation

- Solve the generalized eigenvalue problem for SW and SB.
- Parallelization: Parallelize the eigenvalue and eigenvector computation.

AREAS FOR PARALLELIZATION – LDA 3


7. Projection onto Fisherfaces

- Project the training and test data onto the Fisherfaces.
- Parallelization: Process each image independently during projection.


8. Reconstruction

- Reconstruct images from their projections.
- Parallelization: Distribute reconstruction computations across threads or CUDA blocks.


**OpenMP LDA: Total Execution Time vs Thread Count
(400 Images)**




OpenMP LDA: Parallel Operation Scaling




OpenMP LDA: Total Execution Time (8 Threads)




OpenMP LDA: LDA-Specific Operations




OpenMP LDA: Project Data Performance




CUDA LDA: Total Execution Time




CUDA LDA: Processing Throughput




CUDA LDA: Kernel Performance



LDA Implementation Comparison: Total Execution Time



LDA Implementation Throughput Comparison



LDA Implementation Speedup Comparison

